

```
# Deep learning packages
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torch.utils.data as data_utils
import torchvision.transforms as transforms
from torchvision import datasets
```

```
# Show the shape of the dataset
print("Image shape: {}".format(img.shape))
print("Labels shape: {}".format(labels.shape))
```

```
Image shape: (42000, 784)
Labels shape: (42000,)
```

```
df.head()
```

	label	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	...	pixel774	pixel775	pixel776	pixel777	p
0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
2	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
3	4.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	

5 rows x 785 columns



```
df["label"].value_counts(ascending=True)
```

```
5.0    3795
8.0    4063
4.0    4072
0.0    4132
6.0    4137
2.0    4177
9.0    4188
3.0    4351
7.0    4401
1.0    4684
Name: label, dtype: int64
```

```
df["label"].value_counts(normalize=True)*100
```

```
1.0    11.152381
7.0    10.478571
3.0    10.359524
9.0    9.971429
2.0    9.945238
6.0    9.850000
0.0    9.838095
4.0    9.695238
8.0    9.673810
5.0    9.035714
Name: label, dtype: float64
```

```
img = img.reshape(-1, 1, 28, 28)
```

[illegible]

```

                                test_size=0.25,
                                random_state=42,
                                stratify=labels)

# Split validation/test set into validation and test set
img_val, img_test, labels_val, labels_test = train_test_split(img_val_test,
                                                                labels_val_test,
                                                                test_size=0.5,
                                                                random_state=42,
                                                                stratify=labels_val_test)

```

```

# Define batch_size, epoch and iteration
batch_size = 100
n_iters = 2000
num_epochs = n_iters / (len(img_train) / batch_size)
num_epochs = int(num_epochs)

#-----

# Convert train set to tensors
img_train = torch.from_numpy(img_train)
labels_train = torch.from_numpy(labels_train).type(torch.LongTensor)

# Convert validation set to tensors
img_val = torch.from_numpy(img_val)
labels_val = torch.from_numpy(labels_val).type(torch.LongTensor)

# Convert test set to tensors
img_test = torch.from_numpy(img_test)
labels_test = torch.from_numpy(labels_test).type(torch.LongTensor)

#-----

# Define Pytorch train and validation set
train = data_utils.TensorDataset(img_train, labels_train)
val = data_utils.TensorDataset(img_val, labels_val)
test = data_utils.TensorDataset(img_test, labels_test)

#-----

# Define data loader
train_loader = data_utils.DataLoader(train,
                                     batch_size=batch_size,
                                     shuffle=True, num_workers=16)

valid_loader = data_utils.DataLoader(val,
                                     batch_size=batch_size,
                                     shuffle=True, num_workers=16)

test_loader = data_utils.DataLoader(test,
                                    batch_size=batch_size,
                                    shuffle=True, num_workers=16)

loaders = {'train': train_loader,
           'valid': valid_loader,
           'test': test_loader}

```

/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:561: UserWarning: This DataLoader will create 16 worker processes with `num_workers=16` argument. This will only work on machines with 16 or more CPUs/GPUs. Beware that this will significantly increase your memory footprint.

CNN Architecture

```

num_classes = 10

# check if CUDA is available
use_cuda = torch.cuda.is_available()

# Define the CNN architecture
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN
        self.conv1 = nn.Conv2d(in_channels=1, out_channels=32, kernel_size=5, padding=1)
        self.conv2 = nn.Conv2d(in_channels=32, out_channels=32, kernel_size=5, padding=1)
        self.conv3 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, padding=1)
        self.conv4 = nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3, padding=1)
        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=1)
        self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.Linear(64 * 11 * 11, 2048)
        self.fc2 = nn.Linear(2048, num_classes)
        self.dropout = nn.Dropout(0.5)

```

```

def forward(self, x):
    ## Define forward behavior
    x = self.conv1(x)
    x = self.pool1(F.relu(self.conv2(x)))
    x = self.dropout(x)
    x = self.conv3(x)
    x = self.pool2(F.relu(self.conv4(x)))
    x = self.dropout(x)
    #print(x.shape)
    x = x.view(-1, 64 * 11 * 11)
    x = self.dropout(x)
    x = F.relu(self.fc1(x))
    x = self.dropout(x)
    x = self.fc2(x)
    return x

# instantiate the CNN
model = Net()

# move tensors to GPU if CUDA is available
if use_cuda:
    model.cuda()

print(model)

Net(
  (conv1): Conv2d(1, 32, kernel_size=(5, 5), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(32, 32, kernel_size=(5, 5), stride=(1, 1), padding=(1, 1))
  (conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv4): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (pool1): MaxPool2d(kernel_size=2, stride=1, padding=0, dilation=1, ceil_mode=False)
  (pool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): Linear(in_features=7744, out_features=2048, bias=True)
  (fc2): Linear(in_features=2048, out_features=10, bias=True)
  (dropout): Dropout(p=0.5, inplace=False)
)

```

```

### Define loss function
criterion = nn.CrossEntropyLoss()

### Define optimizer
optimizer = optim.RMSprop(model.parameters(),
                           lr=0.001,
                           alpha=0.9,
                           eps=1e-08,
                           weight_decay=0.0001)

```

```

# The following import is required for training to be robust to truncated images
from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf
    loss_list = []
    epoch_list = []

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
        #####
        model.train()
        for batch_idx, (data, target) in enumerate(loaders['train']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            ## find the loss and update the model parameters accordingly
            optimizer.zero_grad()
            output = model(data)
            loss = criterion(output, target)
            loss.backward()
            optimizer.step()
            train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

        #####
        # validate the model #

```

```
#####
model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    ## update the average validation loss
    output = model(data)
    loss = criterion(output, target)
    valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

## Save the model if validation loss has decreased
if valid_loss < valid_loss_min:
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(valid_loss_min, valid_loss))
    torch.save(model.state_dict(), save_path)
    valid_loss_min = valid_loss

# return trained model
return model
```

```
# Train the model
model = train(10, loaders, model, optimizer,
              criterion, use_cuda, 'cnn_digit_recognizer.pt')
```

```
Epoch: 1      Training Loss: 0.381951      Validation Loss: 0.110209
Validation loss decreased (inf --> 0.110209). Saving model ...
Epoch: 2      Training Loss: 0.133255      Validation Loss: 0.079129
Validation loss decreased (0.110209 --> 0.079129). Saving model ...
Epoch: 3      Training Loss: 0.120103      Validation Loss: 0.066536
Validation loss decreased (0.079129 --> 0.066536). Saving model ...
Epoch: 4      Training Loss: 0.108853      Validation Loss: 0.054784
Validation loss decreased (0.066536 --> 0.054784). Saving model ...
Epoch: 5      Training Loss: 0.108210      Validation Loss: 0.070331
Epoch: 6      Training Loss: 0.105188      Validation Loss: 0.055226
Epoch: 7      Training Loss: 0.102114      Validation Loss: 0.059520
Epoch: 8      Training Loss: 0.102387      Validation Loss: 0.085768
Epoch: 9      Training Loss: 0.102462      Validation Loss: 0.049350
Validation loss decreased (0.054784 --> 0.049350). Saving model ...
Epoch: 10     Training Loss: 0.101559      Validation Loss: 0.080661
```

```
def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

# call test function
test(loaders, model, criterion, use_cuda)
```

```
Test Loss: 0.087433
```

Test Accuracy: 97% (5127/5250)

```
# helper function to un-normalize and display an image
def imshow(img):
    img = img.numpy() * 255 # unnormalize and convert from Tensor image
    plt.imshow(img[0]) # show image

classes = ['0', '1', '2', '3', '4',
           '5', '6', '7', '8', '9']

# obtain one batch of test images
dataiter = iter(test_loader)
images, labels = next(dataiter)
images.numpy()

# move model inputs to cuda, if GPU available
if use_cuda:
    images = images.cuda()

# get sample outputs
output = model(images)
# convert output probabilities to predicted class
_, preds_tensor = torch.max(output, 1)
preds = np.squeeze(preds_tensor.numpy()) if not use_cuda else np.squeeze(preds_tensor.cpu().numpy())

# plot the images in the batch, along with predicted and true labels
fig = plt.figure(figsize=(20, 5))
for idx in np.arange(10):
    ax = fig.add_subplot(2, 5, idx+1, xticks=[], yticks=[])
    imshow(images.cpu()[idx])
    ax.set_title("Predicted: {} - Actual: {}".format(classes[preds[idx]], classes[labels[idx]]),
                color=("green" if preds[idx]==labels[idx].item() else "red"))
```

